

Aggregations with pandas and numpy About the Data In this notebook, we will be working with 2 data sets:

Facebook's stock price throughout 2018 (obtained using the stock_analysis package). daily weather data for NYC from the National Centers for Environmental Information (NCEI) API. Note: The NCEI is part of the National Oceanic and Atmospheric Administration (NOAA) and, as you can see from the URL for the API, this resource was created when the NCEI was called the NCDC. Should the URL for this resource change in the future, you can search for the NCEI weather API to find the updated one.

Background on the weather data Data meanings:

AWND : average wind speed PRCP : precipitation in millimeters SNOW : snowfall in millimeters - SNWD : snow depth in millimeters TMAX : maximum daily temperature in Celsius TMIN : minimum daily temperature in Celsius

```
import numpy as np
import pandas as pd
```

```
weather = pd.read_csv('weather_by_station.csv', index_col='date', parse_dates=True)
weather.head()
```



	datatype	station	value	station_name
date				
2018-01-01	PRCP	GHCND:US1CTFR0039	0.0	STAMFORD 4.2 S, CT US
2018-01-01	PRCP	GHCND:US1NJBG0015	0.0	NORTH ARLINGTON 0.7 WNW, NJ US
2018-01-01	SNOW	GHCND:US1NJBG0015	0.0	NORTH ARLINGTON 0.7 WNW, NJ US
2018-01-				



Next steps: [View recommended plots](#) [New interactive sheet](#)

```
fb = pd.read_csv('/content/fb_2018 8.3.csv', index_col='date', parse_dates=True).assign(
    trading_volume=lambda x: pd.cut(x.volume, bins=3, labels=['low', 'med', 'high']))
fb.head()
```



	open	high	low	close	volume	trading_volume	
date							
2018-01-02	177.68	181.58	177.5500	181.42	18151903	low	
2018-01-03	181.88	184.78	181.3300	184.67	16886563	low	
2018-01-04	184.90	186.21	184.0996	184.33	13880896	low	
2018-01-05	185.59	186.90	184.9300	186.85	13574535	low	
2018-01-08	187.20	188.90	186.3300	188.28	17994726	low	

Next steps:

[View recommended plots](#)[New interactive sheet](#)

```
pd.set_option('display.float_format', lambda x: '%.2f' % x)
```

Summarizing DataFrames

We learned about `agg()` in the dataframe operations notebook when we learned about window calculations; however, we can call this on the dataframe directly to aggregate its contents into a single series:

```
fb.agg({
    'open': 'mean',
    'high': 'max',
    'low': 'min',
    'close': 'mean',
    'volume': 'sum'
})
```



```

      0
open    171.45
high    218.62
low     123.02
close    171.51
volume  6949682394.00

dtype: float64
```

```
weather.query(
    'station == "GHCND:USW00094728"'
).pivot(columns='datatype', values='value')[['SNOW', 'PRCP']].sum()
```



0

datatype**SNOW** 1007.00**PRCP** 1665.30**dtype:** float64

```
weather.query(
    'station == "GHCND:USW00094728"'
).pivot(columns='datatype', values='value')[['SNOW', 'PRCP']].agg('sum')
```



0

datatype**SNOW** 1007.00**PRCP** 1665.30**dtype:** float64

```
fb.agg({
    'open': 'mean',
    'high': ['min', 'max'],
    'low': ['min', 'max'],
    'close': 'mean'
})
```




	open	high	low	close
mean	171.45	NaN	NaN	171.51
min	NaN	129.74	123.02	NaN
max	NaN	218.62	214.27	NaN





Using groupby()

Often we won't want to aggregate on the entire dataframe, but on groups within it. For this purpose, we can run `groupby()` before the aggregation. If we group by the `trading_volume` column, we will get a row for each of the values it takes on:


```
fb.groupby('trading_volume', observed=True).mean()
```





	open	high	low	close	volume
trading_volume					
low	171.36	173.46	169.31	171.43	24547207.71
med	175.82	179.42	172.11	175.14	79072559.12
high	167.73	170.48	161.57	168.16	141924023.33




```
fb.groupby('trading_volume',observed=True)['close'].agg(['min', 'max', 'mean'])
```






	min	max	mean
trading_volume			
low	124.06	214.67	171.43
med	152.22	217.50	175.14
high	160.06	176.26	168.16



```
fb_agg = fb.groupby('trading_volume', observed = True).agg({
    'open': 'mean',
    'high': ['min', 'max'],
    'low': ['min', 'max'],
    'close': 'mean'
})
fb_agg
```



	open	high		low		close
	mean	min	max	min	max	mean
trading_volume						
low	171.36	129.74	216.20	123.02	212.60	171.43
med	175.82	162.85	218.62	150.75	214.27	175.14
high	167.73	161.10	180.13	149.02	173.75	168.16



Next steps:


[View recommended plots](#)

[New interactive sheet](#)



```
fb_agg.columns
MultiIndex([( 'open', 'mean'),
            ( 'high', 'min'),
            ( 'high', 'max'),
            ( 'low', 'min')])
```

```
( 'low', 'max'),
('close', 'mean')],
)
```

```
fb_agg.columns = ['_'.join(col_agg) for col_agg in fb_agg.columns]
fb_agg.head()
```



	open_mean	high_min	high_max	low_min	low_max	close_mean
trading_volume						
low	171.36	129.74	216.20	123.02	212.60	171.43
med	175.82	162.85	218.62	150.75	214.27	175.14
high	167.73	161.10	180.13	149.02	173.75	168.16




Next steps:

 View recommended plots

New interactive sheet

```
weather.loc['2018-10'].query('datatype == "PRCP"]').groupby(
    pd.Grouper(freq='D')
)['value'].mean().head()
```



	value
date	
2018-10-01	0.01
2018-10-02	2.23
2018-10-03	19.69
2018-10-04	0.32
2018-10-05	0.97

dtype: float64

```
weather.query('datatype == "PRCP"]').groupby(
    ['station_name', pd.Grouper(freq='QE')])
['value'].sum().unstack().sample(5, random_state=1)
```



date 2018-03-31 2018-06-30 2018-09-30 2018-12-31



station_name



WANTAGH 1.1 NNE, NY US	279.90	216.80	472.50	277.20
STATEN ISLAND 1.4 SE, NY US	379.40	295.30	438.80	409.90
SYOSSET 2.0 SSW, NY US	323.50	263.30	355.50	459.90
STAMFORD 4.2 S, CT US	338.00	272.10	424.70	390.00
WAYNE TWP 0.8 SSW, NJ US	246.20	295.30	620.90	422.00

```
weather.groupby('station').filter(
    lambda x: 'NY' in x.name
).query('datatype == "SNOW"]').groupby('station_name')['value'].sum().squeeze()
```



	value
station_name	
ALBERTSON 0.2 SSE, NY US	1087.00
AMITYVILLE 0.1 WSW, NY US	434.00
AMITYVILLE 0.6 NNE, NY US	1072.00
ARMONK 0.3 SE, NY US	1504.00
BROOKLYN 3.1 NW, NY US	305.00
CENTERPORT 0.9 SW, NY US	799.00
ELMSFORD 0.8 SSW, NY US	863.00
FLORAL PARK 0.4 W, NY US	1015.00
HICKSVILLE 1.3 ENE, NY US	716.00
JACKSON HEIGHTS 0.3 WSW, NY US	107.00
LOCUST VALLEY 0.3 E, NY US	0.00
LYNBROOK 0.3 NW, NY US	325.00
MASSAPEQUA 0.9 SSW, NY US	41.00
MIDDLE VILLAGE 0.5 SW, NY US	1249.00
NEW HYDE PARK 1.6 NE, NY US	0.00
NEW YORK 8.8 N, NY US	0.00
NORTH WANTAGH 0.4 WSW, NY US	471.00
PLAINEDGE 0.4 WSW, NY US	610.00
PLAINVIEW 0.4 ENE, NY US	1360.00
SADDLE ROCK 3.4 WSW, NY US	707.00
STATEN ISLAND 1.4 SE, NY US	936.00
STATEN ISLAND 4.5 SSE, NY US	89.00
SYOSSET 2.0 SSW, NY US	1039.00
VALLEY STREAM 0.6 SE, NY US	898.00
WANTAGH 0.3 ESE, NY US	1280.00
WANTAGH 1.1 NNE, NY US	940.00
WEST NYACK 1.3 WSW, NY US	1371.00

dtype: float64

```
weather.query('datatype == "PRCP"]').groupby(
    pd.Grouper(freq='D')
)['value'].mean().groupby(pd.Grouper(freq='ME')).sum().nlargest()
```



	value
date	
2018-11-30	210.59
2018-09-30	193.09
2018-08-31	192.45
2018-07-31	160.98
2018-02-28	158.11

dtype: float64

Perhaps the previous result was surprising. The saying goes "April showers bring May flowers"; yet April wasn't in the top 5 (neither was May for that matter). Snow will count towards precipitation, but that doesn't explain why summer months are higher than April. Let's look for days that accounted for a large percentage of the precipitation in a given month. In order to do so, we need to calculate the average daily precipitation across stations and then find the total per month. This will be the denominator. However, in order to divide the daily values by the total for their month, we will need a Series of equal dimensions. This means we will need to use `transform()` :

```
weather.query('datatype == "PRCP"]').rename(
    columns={'value': 'prcp'})
.groupby(pd.Grouper(freq='D'))['prcp'].mean().groupby(
    pd.Grouper(freq='ME'))
.transform('sum')['2018-01-28': '2018-02-03']
```




	prcp
date	
2018-01-28	69.31
2018-01-29	69.31
2018-01-30	69.31
2018-01-31	69.31
2018-02-01	158.11
2018-02-02	158.11
2018-02-03	158.11

dtype: float64

```
weather \
.query('datatype == "PRCP"') \
.rename(columns={'value': 'prcp'}) \
.groupby(pd.Grouper(freq='D'))['prcp'].mean() \
.to_frame().assign(
    total_prdp_in_month=lambda x: x.groupby(pd.Grouper(freq='ME'))['prcp'].transform('sum'),
    pct_monthly_prdp=lambda x: x['prcp'].div(x.total_prdp_in_month)
) \
.nlargest(5, 'pct_monthly_prdp')
```



	prcp	total_prdp_in_month	pct_monthly_prdp
date			
2018-10-12	34.77	105.63	0.33
2018-01-13	21.66	69.31	0.31
2018-03-02	38.77	137.46	0.28
2018-04-16	39.34	140.57	0.28
2018-04-17	37.30	140.57	0.27



```
fb[['open', 'high', 'low', 'close']].transform(
    lambda x: (x - x.mean()).div(x.std())
).head()
```



	open	high	low	close
date				
2018-01-02	0.32	0.41	0.41	0.50
2018-01-03	0.53	0.57	0.60	0.66
2018-01-04	0.68	0.65	0.74	0.64
2018-01-05	0.72	0.68	0.78	0.77
2018-01-08	0.80	0.79	0.85	0.84



Pivot tables and crosstabs

We saw pivots in before; however, we weren't able to provide any aggregations. With `pivot_table()`, we get the mean by default as the `aggfunc`. In its simplest form, we provide a column to place along the columns:

```
fb.pivot_table(columns='trading_volume',observed = True)
```



trading_volume	low	med	high
close	171.43	175.14	168.16
high	173.46	179.42	170.48
low	169.31	172.11	161.57
open	171.36	175.82	167.73
volume	24547207.71	79072559.12	141924023.33



```
fb.pivot_table(index='trading_volume',observed = True)
```




	close	high	low	open	volume
trading_volume					
low	171.43	173.46	169.31	171.36	24547207.71
med	175.14	179.42	172.11	175.82	79072559.12
high	168.16	170.48	161.57	167.73	141924023.33




```
weather.reset_index().pivot_table(
index=['date', 'station', 'station_name']
columns='datatype',
```

```
values='value',
aggfunc='median'
).reset_index().tail()
```



datatype	date	statio	station_name	AWND	DAPR	MDPR	PGTM	PRCP	5
28740	2018-12-31	GHCND:USW00054787	FARMINGDALE REPUBLIC AIRPORT, NY US	5.00	NaN	NaN	2052.00	28.70	
28741	2018-12-31	GHCND:USW00094728	NY CITY CENTRAL PARK, NY US	NaN	NaN	NaN	NaN	25.90	
28742	2018-12-31	GHCND:USW00094741	TETERBORO AIRPORT, NJ US	1.70	NaN	NaN	1954.00	29.20	
28743	2018-12-31	GHCND:USW00094745	WESTCHESTER CO AIRPORT, NY US	2.70	NaN	NaN	2212.00	24.40	



```
pd.crosstab(
index=fb.trading_volume,
columns=fb.index.month,
colnames=['month']
)
```

month	1	2	3	4	5	6	7	8	9	10	11	12	
trading_volume													
low	20	19	15	20	22	21	18	23	19	23	21	19	
med	1	0	4	1	0	0	2	0	0	0	0	0	
high	0	0	2	0	0	0	1	0	0	0	0	0	

```
pd.crosstab(
index=fb.trading_volume,
columns=fb.index.month,
```